

A Hybrid Operating System Design Approach using Hardware Acceleration

Insop Song
Ericsson

Real Time Linux Workshop 2012
UNC Chapel Hill, NC



Hardware OS

- less overhead
- more deterministic

Hardware OS

- less overhead
- more deterministic

Why Hardware OS
NOT seen much?

Hardware OS

- less overhead
- more deterministic

Why Hardware OS
NOT seen much?


Hybrid OS

OS: Process/Timer/Event handling

Process scheduling

OS: Process/Timer/Event handling

Process scheduling

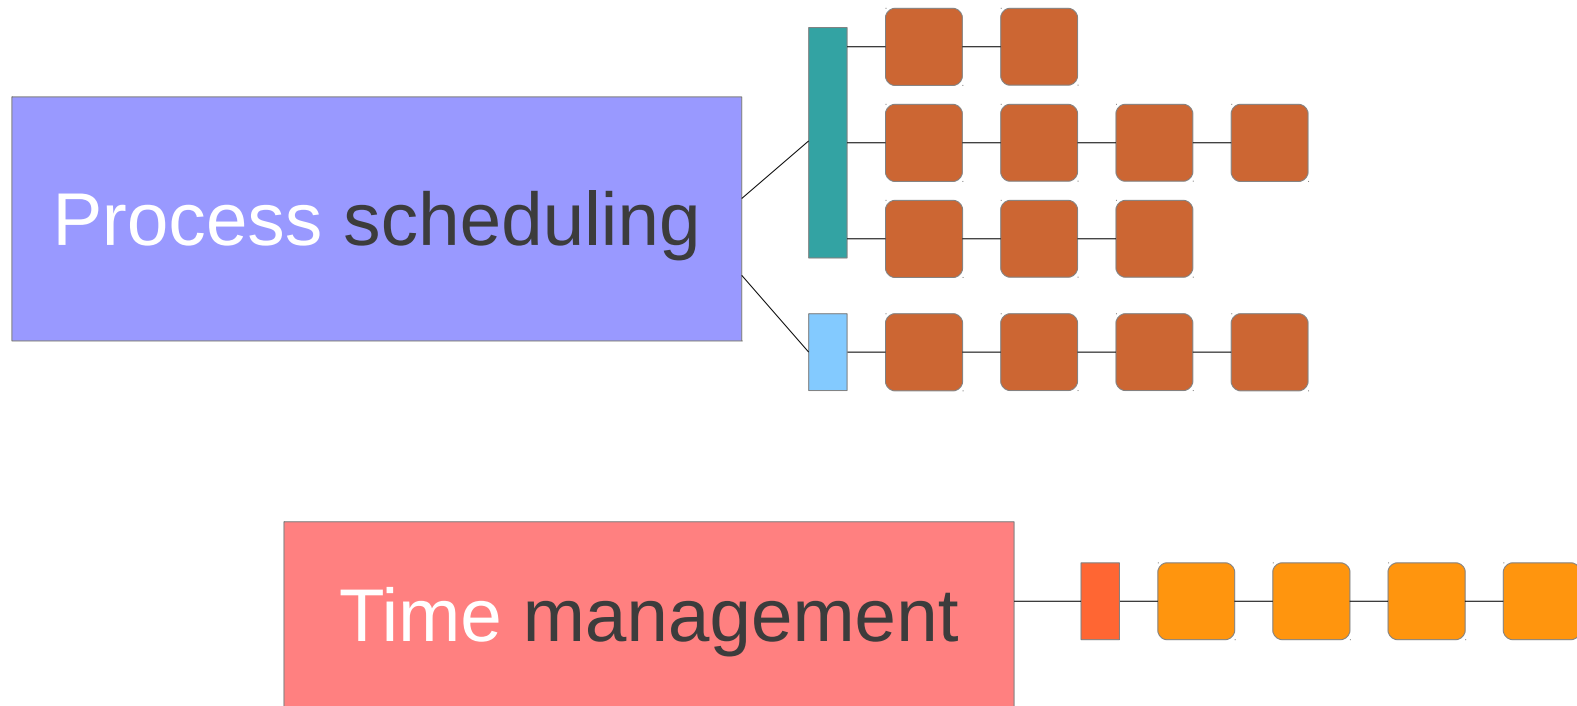


OS: Process/Timer/Event handling

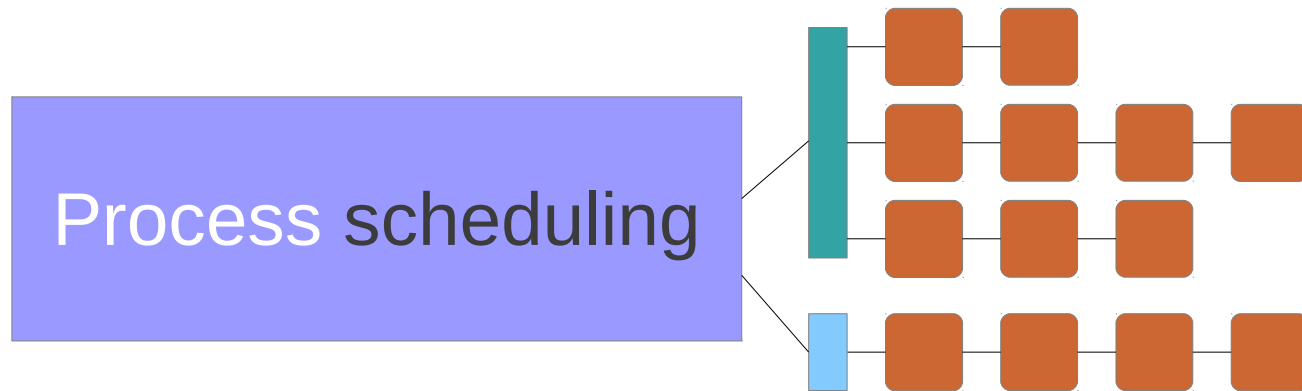
Process scheduling



OS: Process/Timer/Event handling



OS: Process/Timer/Event handling



H/W accelerators & co-processors

- CPU bound
- Process intensive
- Real h/w concurrency

H/W accelerators & co-processors

- CPU bound
- Process intensive
- Real h/w concurrency

Math
coprocessors

H/W accelerators & co-processors

- CPU bound
- Process intensive
- Real h/w concurrency

Math
coprocessors

GPU

H/W accelerators & co-processors

- CPU bound
- Process intensive
- Real h/w concurrency

Math
coprocessors

GPU

DSP

H/W accelerators & co-processors

- CPU bound
- Process intensive
- Real h/w concurrency

Math
coprocessors

GPU

DSP

Video
processing

H/W accelerators & co-processors

- CPU bound
- Process intensive
- Real h/w concurrency

Math
coprocessors

GPU

DSP

Video
processing

Encryption

H/W accelerators & co-processors

- CPU bound
- Process intensive
- Real h/w concurrency

Math
coprocessors

GPU

DSP

Video
processing

FPGA
Custom IP

Encryption

H/W accelerators & co-processors

- CPU bound
- Process intensive
- Real h/w concurrency

Math
coprocessors

GPU

DSP

Video
processing

FPGA
Custom IP

REGEX
accelerator


Encryption

What is Hardware OS?

- Implement Process/Timer/Event handling in H/W
- Not a new idea
 - papers can be found from 90s
- Reduce scheduling overhead
- Improve determinism

What is **Hardware OS**?


- Implement **task/event/timer** handling in H/W
- Not a new idea
 - papers can be found from 90s
- Reduce scheduling **overhead**
- Improve **determinism**




What is **Hardware** OS?

- Implement task/event/timer handling in H/W
- Not a new idea
 - papers can be found from 90s
- Reduce scheduling overhead
- Improve determinism

- Process handling
- Timer handling
- Event handling



H/W blocks inside Hardware OS



Selected List of H/W OS

- Silicon TRON
- F-Timer
- Cooling's scheduler co-processor board
- FASTCHART/RTM (real time unit)
- A. Morton's scheduler coprocessor
- Sierra OS (FPGA IP)

Selected List of H/W OS

- Silicon TRON
- F-Timer
- Cooling's scheduler co-processor board
- FASTCHART/RTM (real time unit)
- A. Morton's scheduler coprocessor
- Sierra OS (FPGA IP)

Selected List of H/W OS

- Silicon TRON
- F-Timer
- Cooling's scheduler co-processor board
- FASTCHART/RTM (real time unit)
- A. Morton's scheduler coprocessor
- Sierra OS (FPGA IP)

Selected List of H/W OS

- Silicon TRON
- F-Timer
- Cooling's scheduler co-processor board
- FASTCHART/RTM (real time unit)
- A. Morton's scheduler coprocessor
- Sierra OS (FPGA IP)

Selected List of H/W OS

- Silicon TRON
- F-Timer
- Cooling's scheduler co-processor board
- FASTCHART/RTM (real time unit)
- A. Morton's scheduler coprocessor
- Sierra OS (FPGA IP)

Selected List of H/W OS

- Silicon TRON
- F-Timer
- Cooling's scheduler co-processor board
- FASTCHART/RTM (real time unit)
- A. Morton's scheduler coprocessor
- Sierra OS (FPGA IP)

Great! H/W OS is available?


- Not quite
- Before we discuss why not
 - Let's talk briefly Linux and Open source software
- Why not h/w OS is available?
 - Need software OS change
 - OS scheduler is quite complex and changing

Linux scheduler and hrtimer

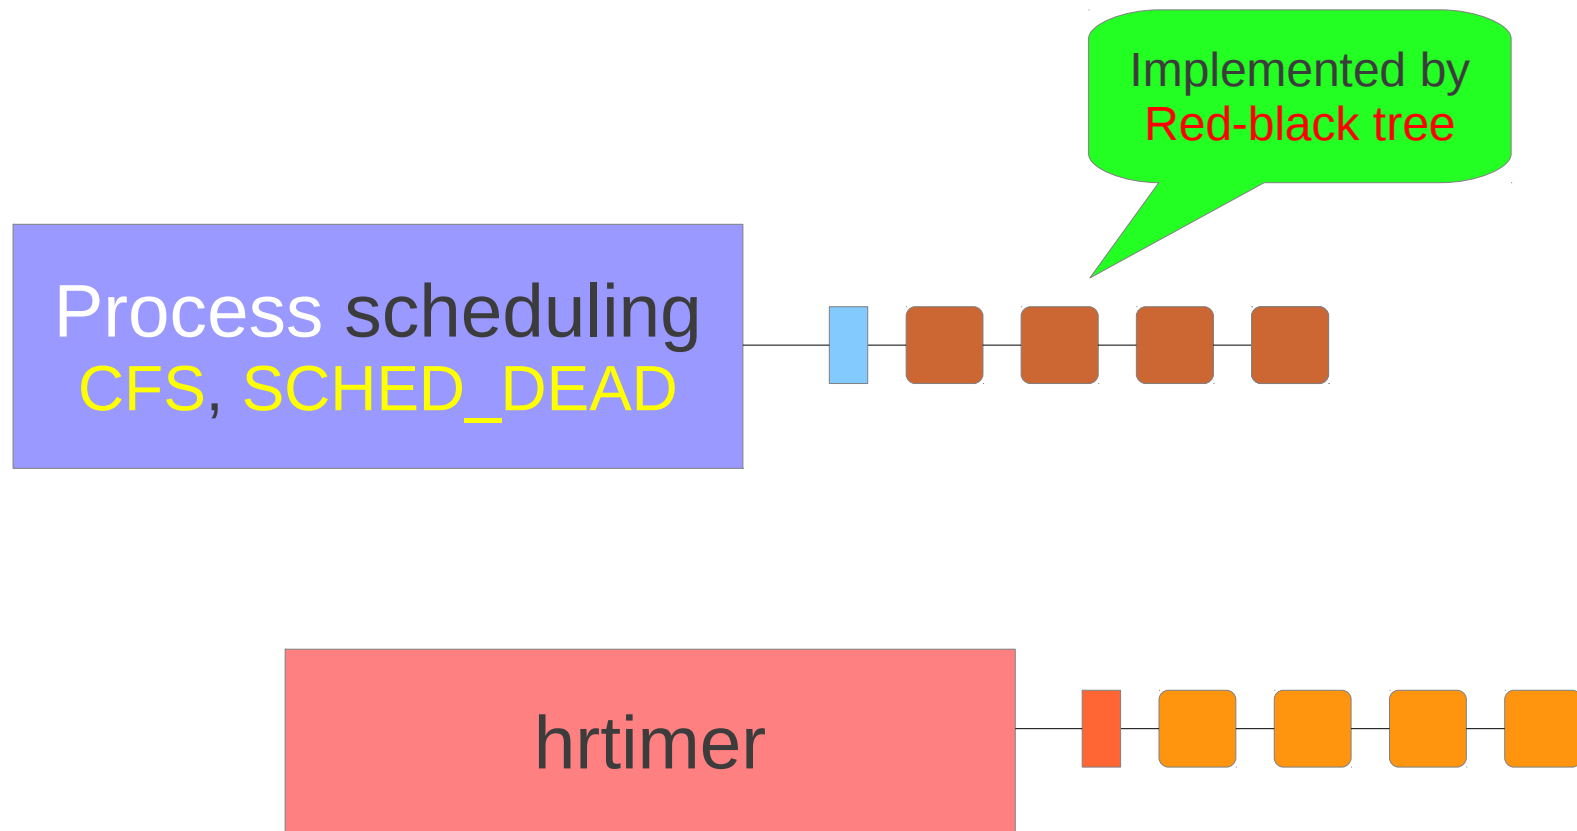
Process scheduling
CFS, SCHED_DEAD



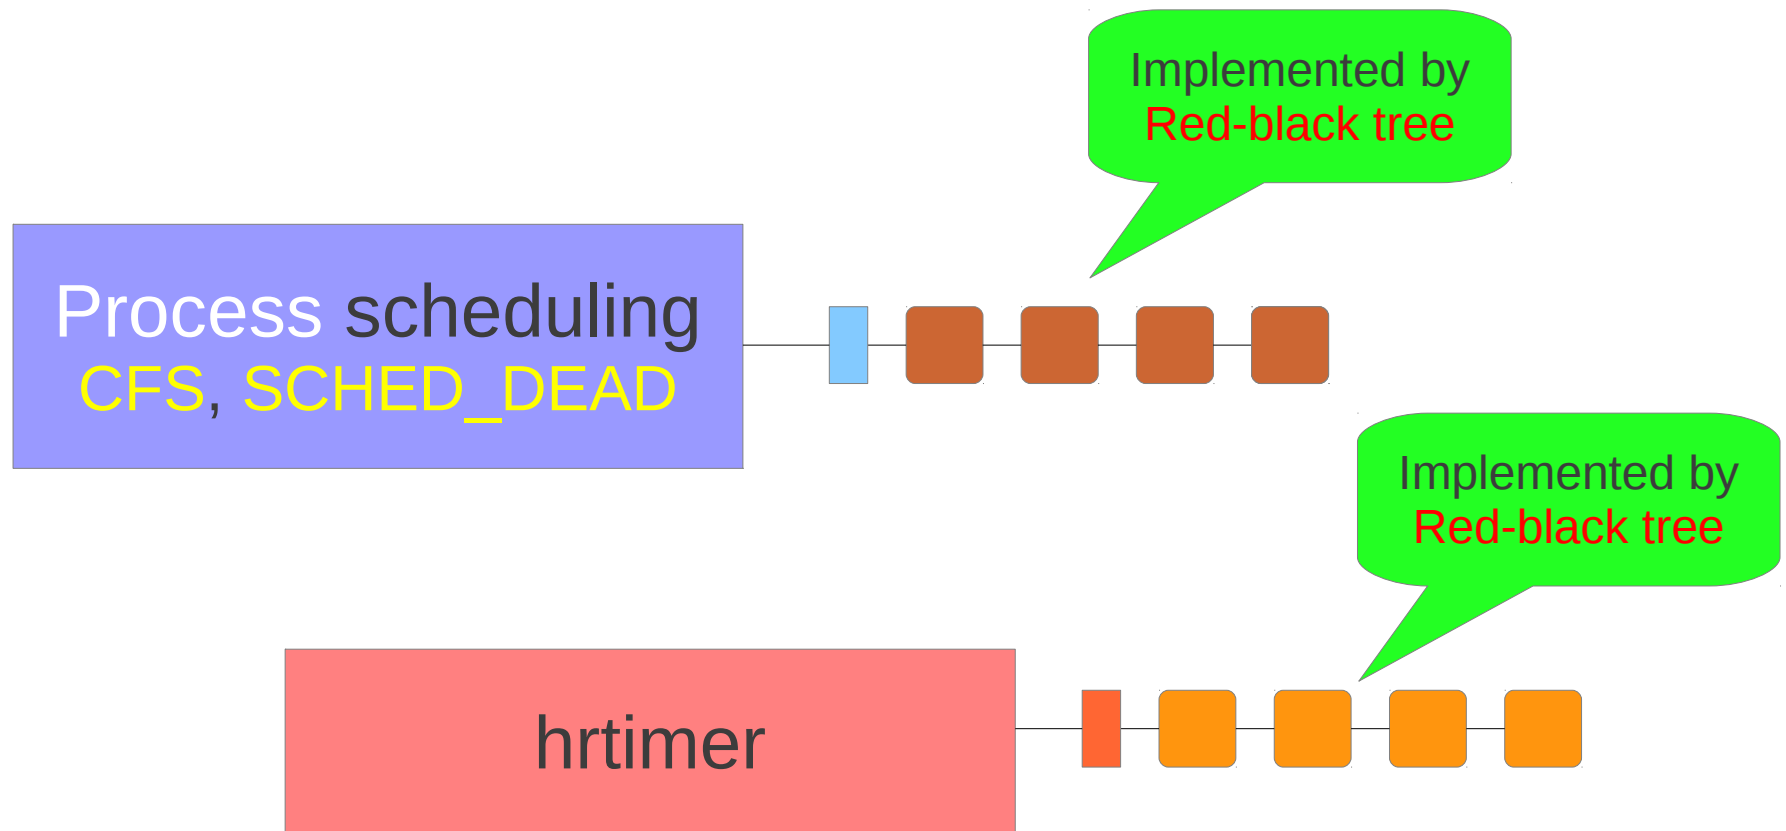
Linux scheduler and hrtimer



Linux scheduler and hrtimer



Linux scheduler and hrtimer




Hybrid approach

- Instead of trying to implement full/semi-full scheduler in h/w
 - Only implement a **priority queue** in h/w
 - Scheduler use **h/w queue** instead of RB-tree


Hybrid approach

- Implement priority queue in h/w
- Scheduler use h/w queue instead of RB-tree



Hybrid approach

- Implement priority queue in h/w
- Scheduler use h/w queue instead of RB-tree



H/W Priority Queue

- Requirements:
 - Enqueue/dequeue in one cycle
 - Maintain FIFO order
- Example application of h/w priority queue
 - Network switch/router traffic management
 - for maintaining QoS during congestion

Types of h/w priority queues

- N: number entry
- P: number priorities

Types of h/w priority queues

- N: number entry
- P: number priorities

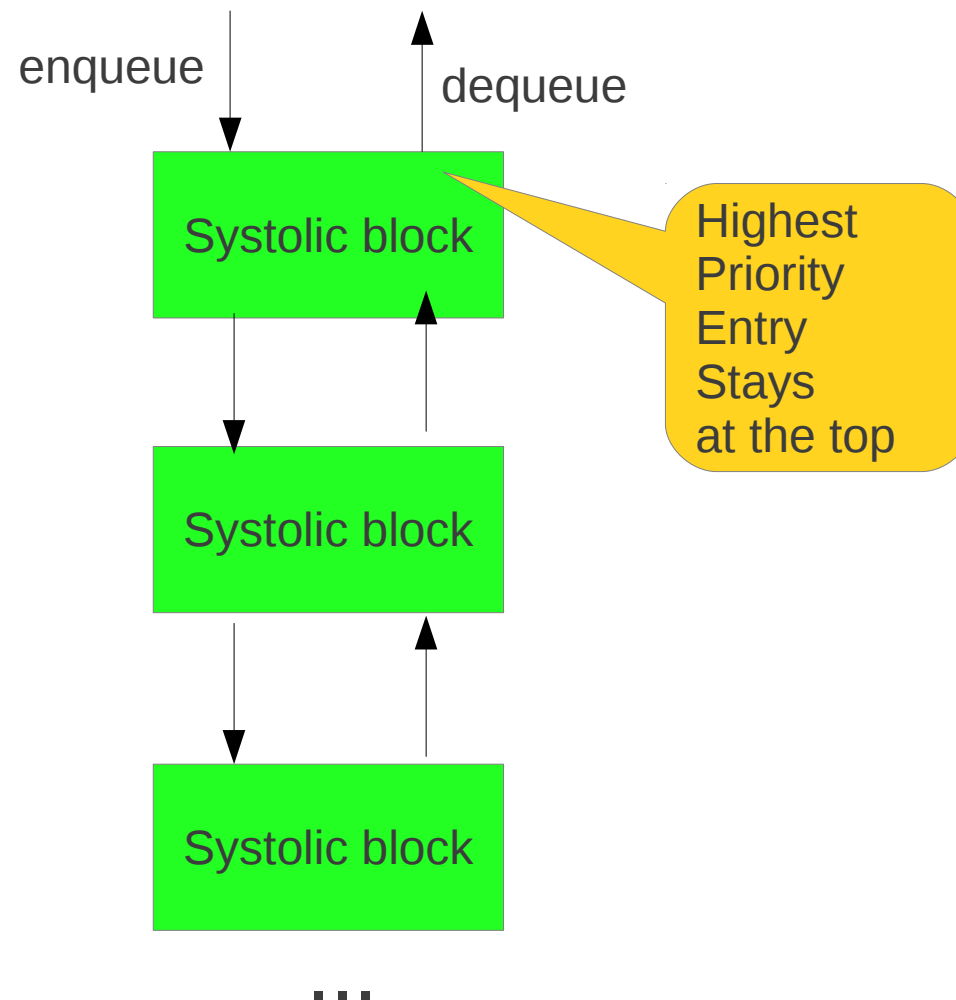
- Binary tree of comparators
- Priority encoder with multiple queues
- Tagged up/down sorter
 - Scales well with large N and P

Types of h/w priority queues

- N: number entry
- P: number priorities


- Binary tree of comparators
- Priority encoder with multiple queues
- Tagged up/down sorter
 - Scales well with large N and P

Tagged Up/Down H/W Sorter




Implementation

- FPGA




Implementation

- FPGA
- Soft-core processor



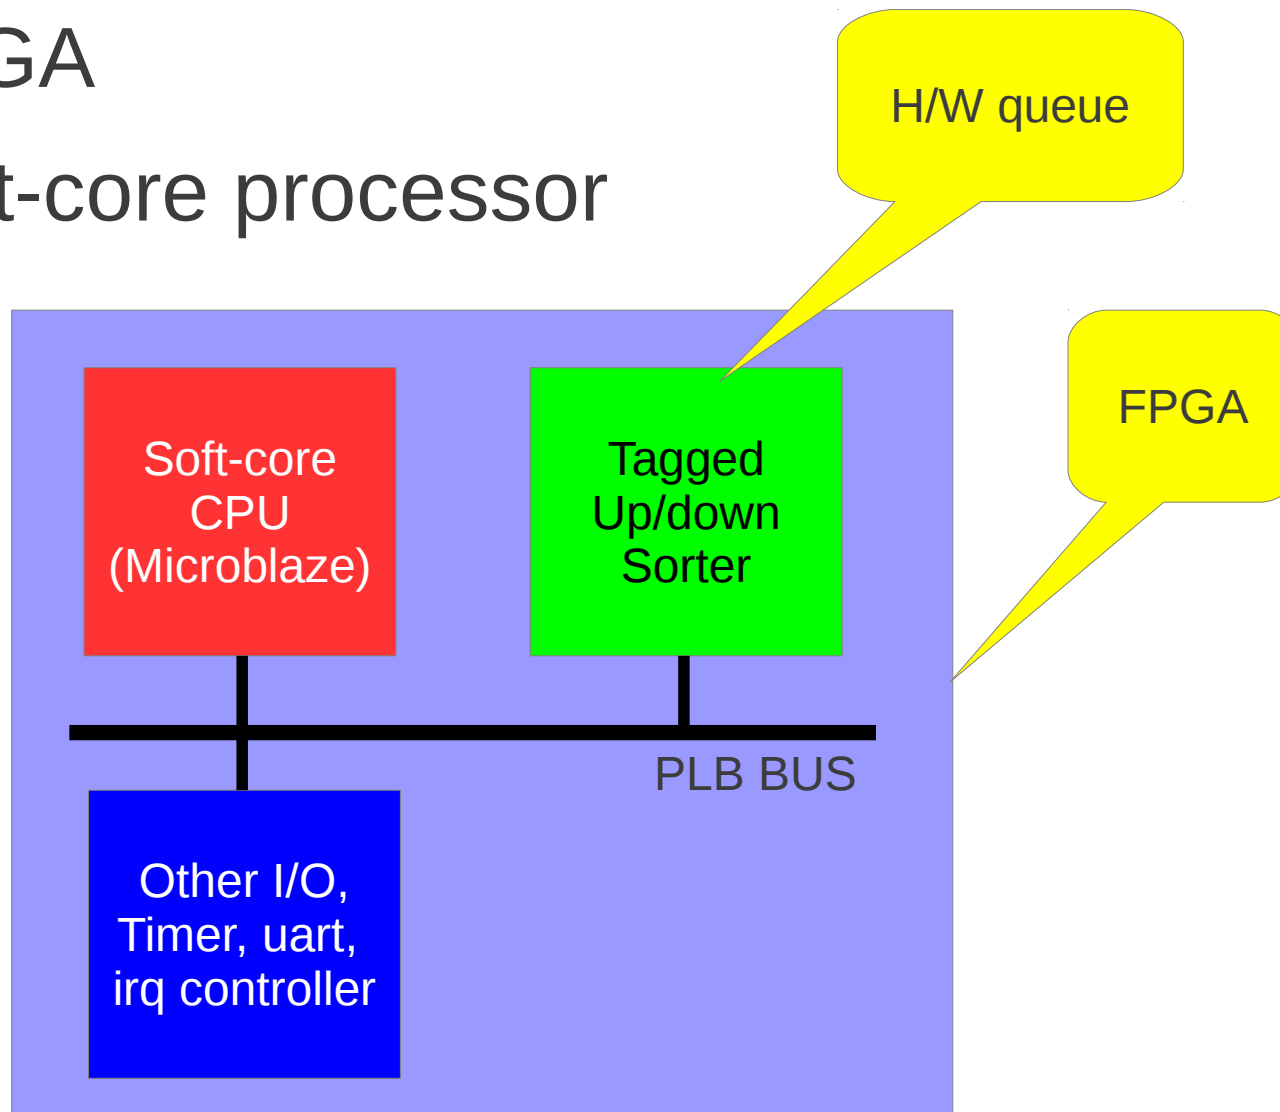
Implementation

- FPGA
- Soft-core processor



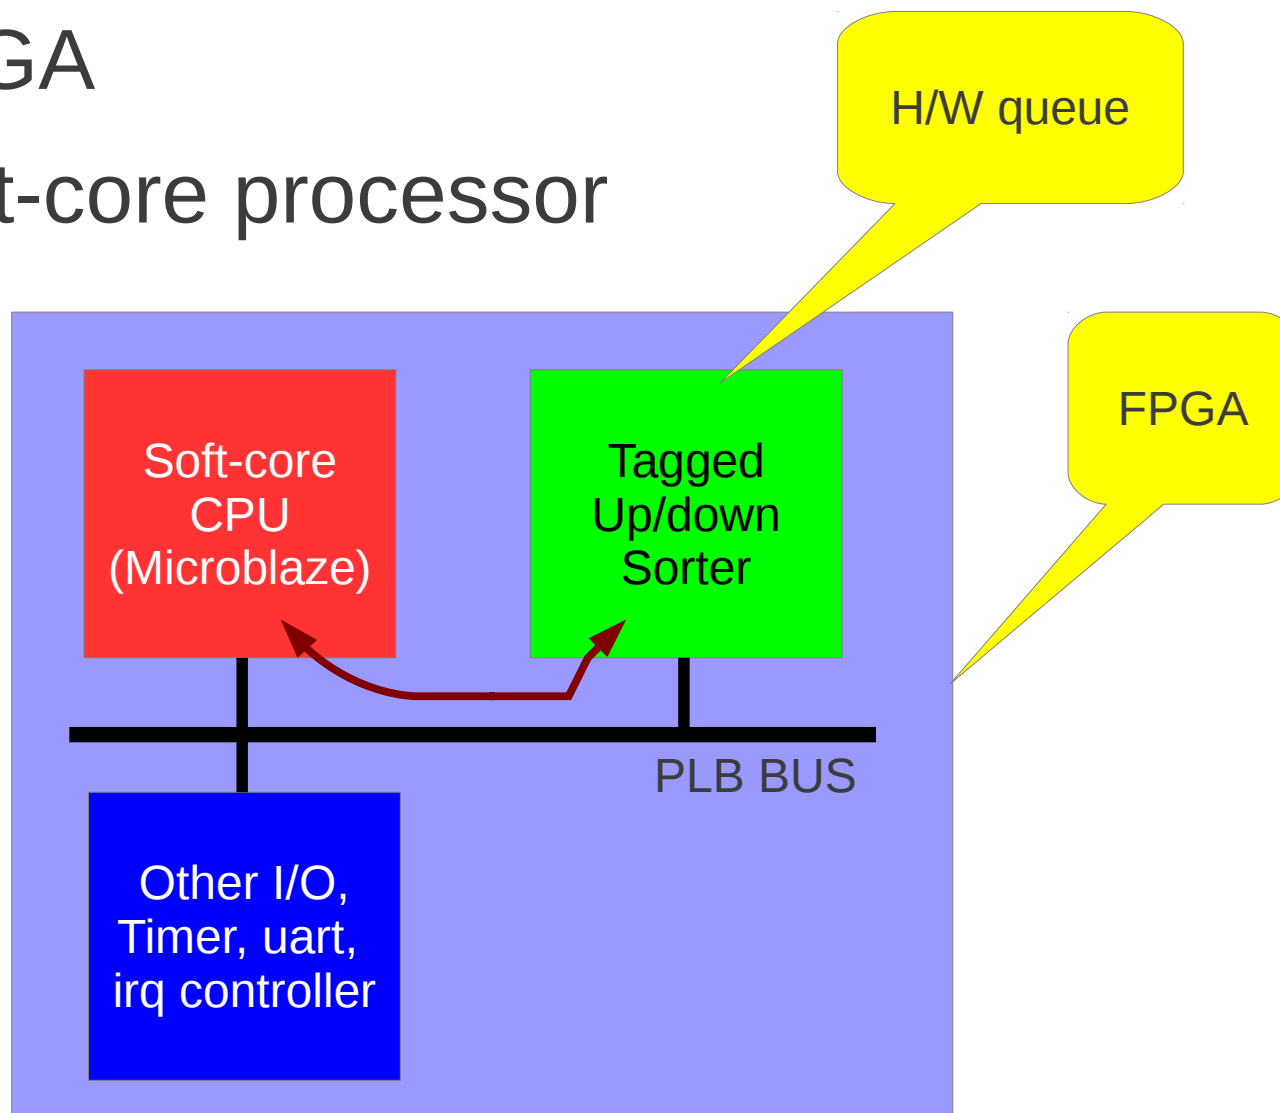
Implementation

- FPGA
- Soft-core processor



Implementation

- FPGA
- Soft-core processor




Experimental setup

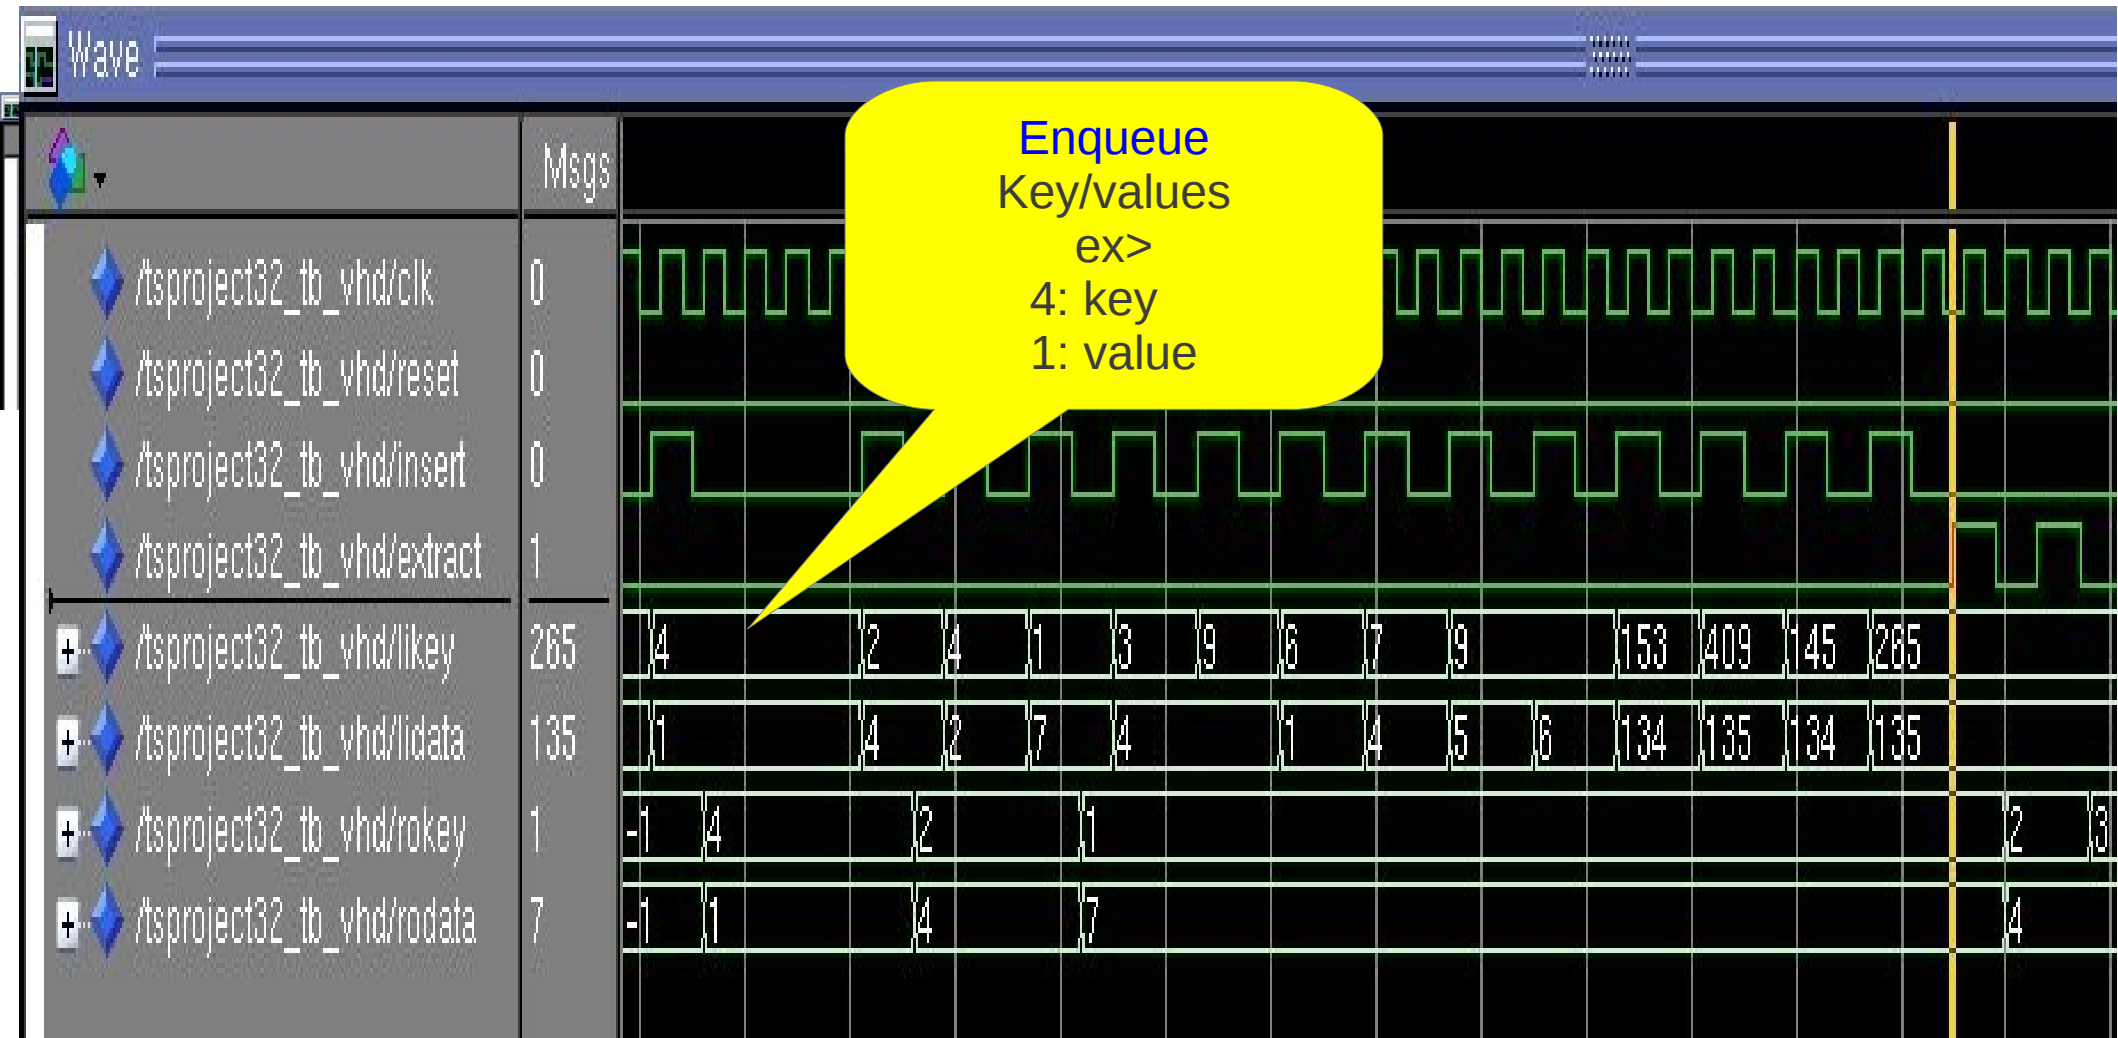
- FPGA (Field Programmable Gate Array)
 - Xilinx Virtex5
- Soft-core CPU
 - Microblaze
- ML505 evaluation board



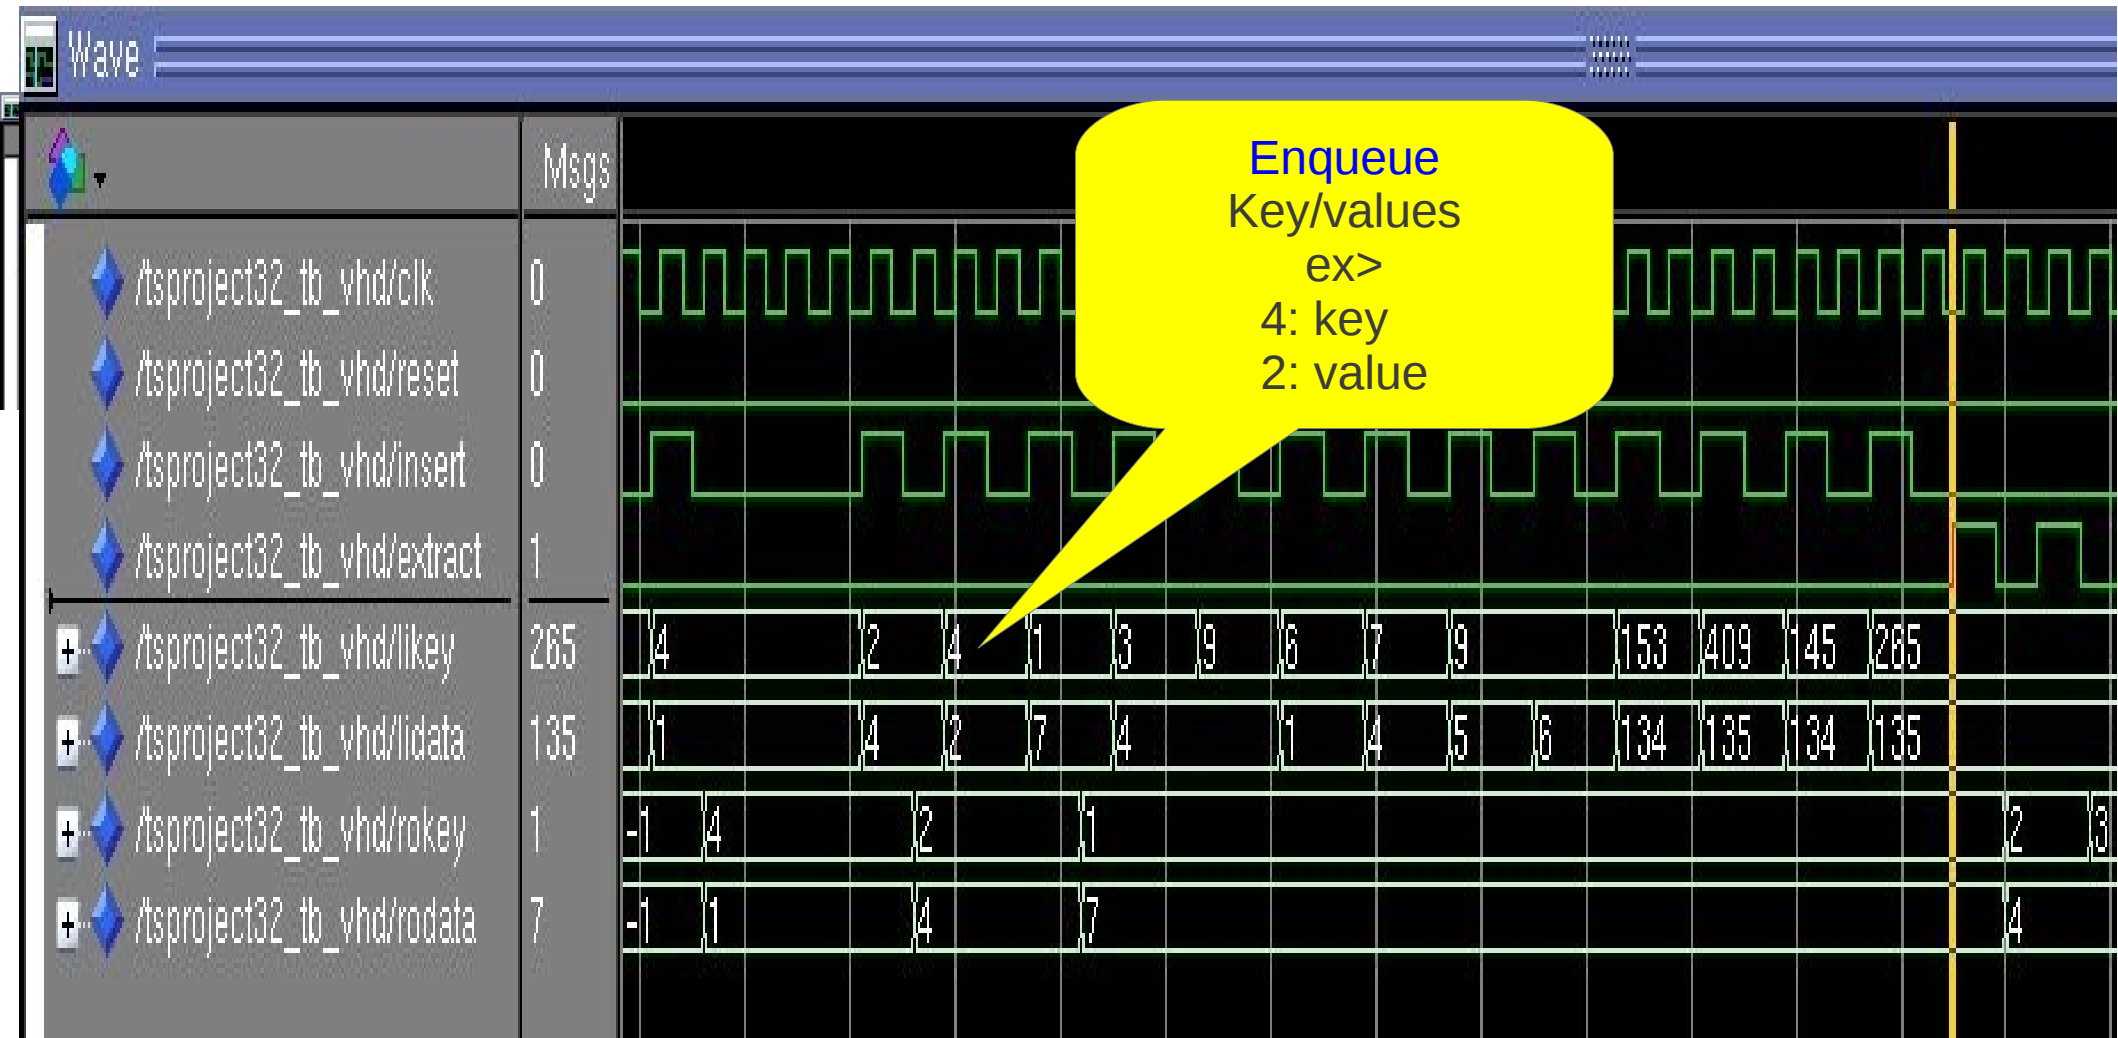
H/W Priority Queue Simulation

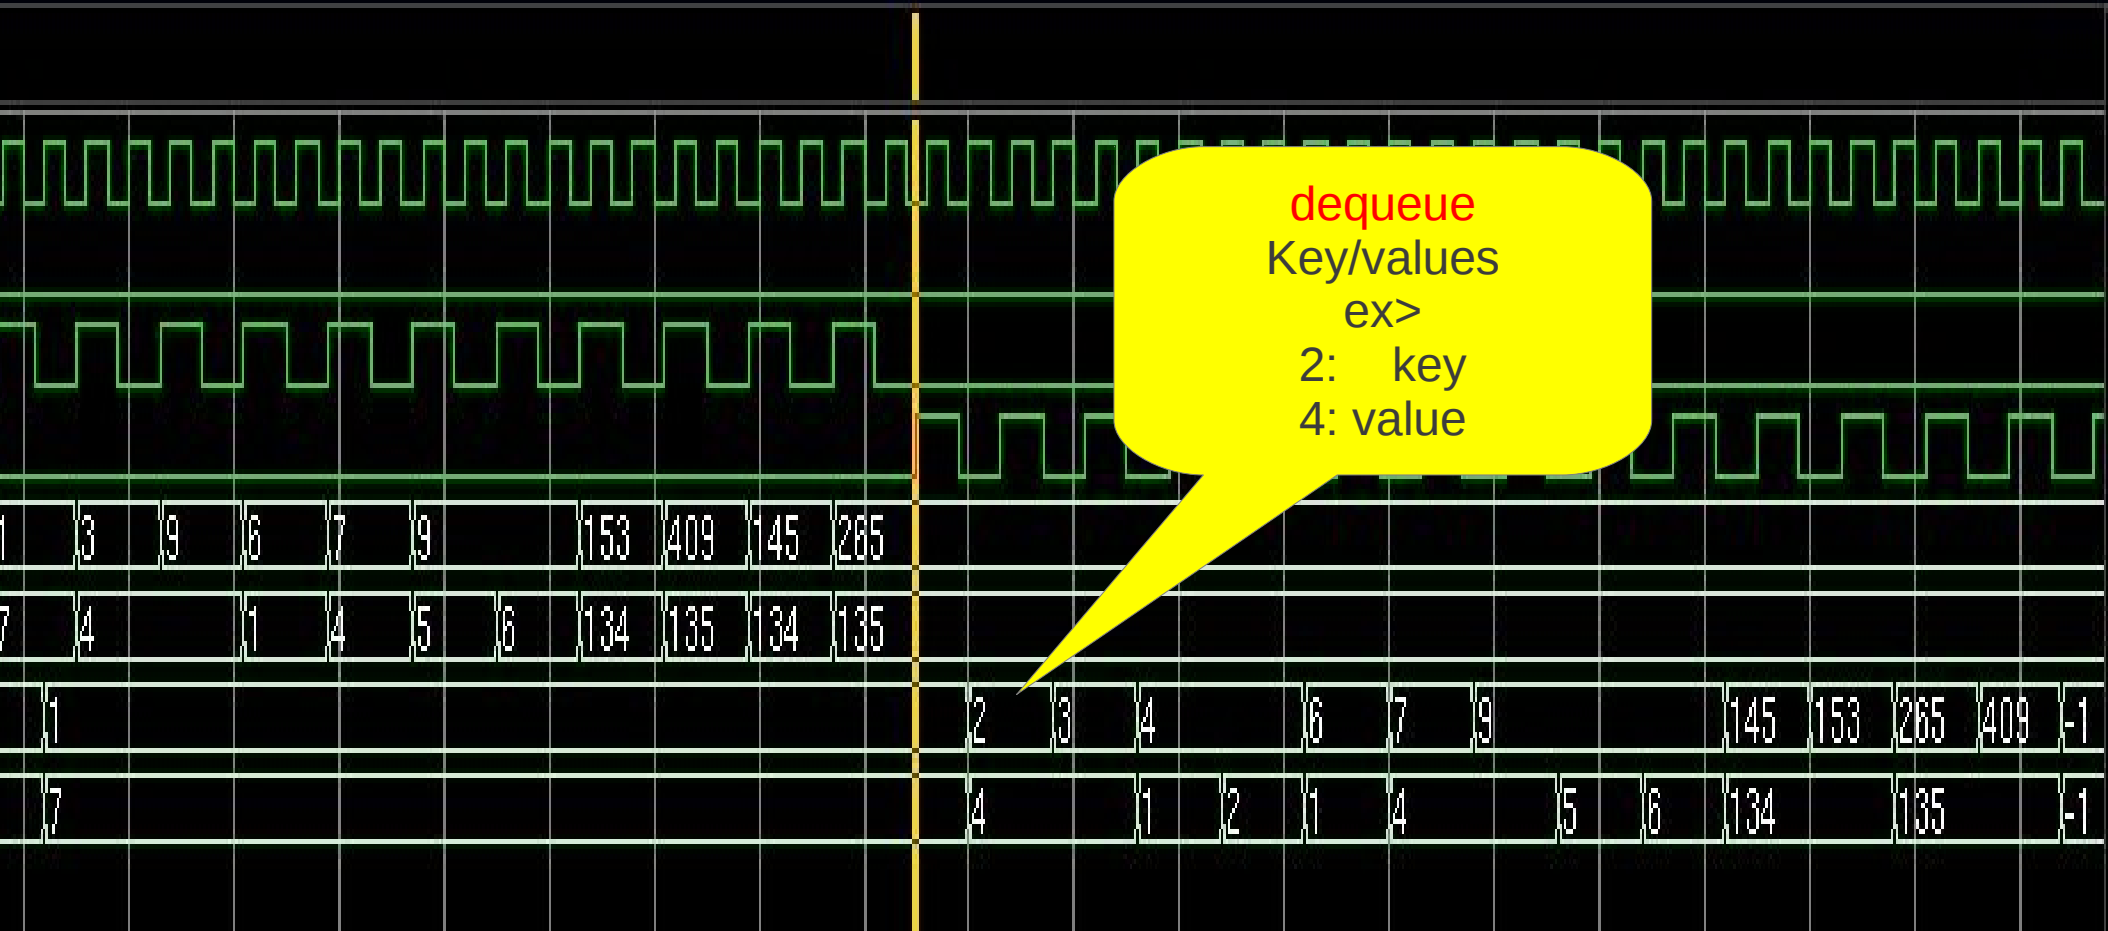


H/W Priority Queue Simulation

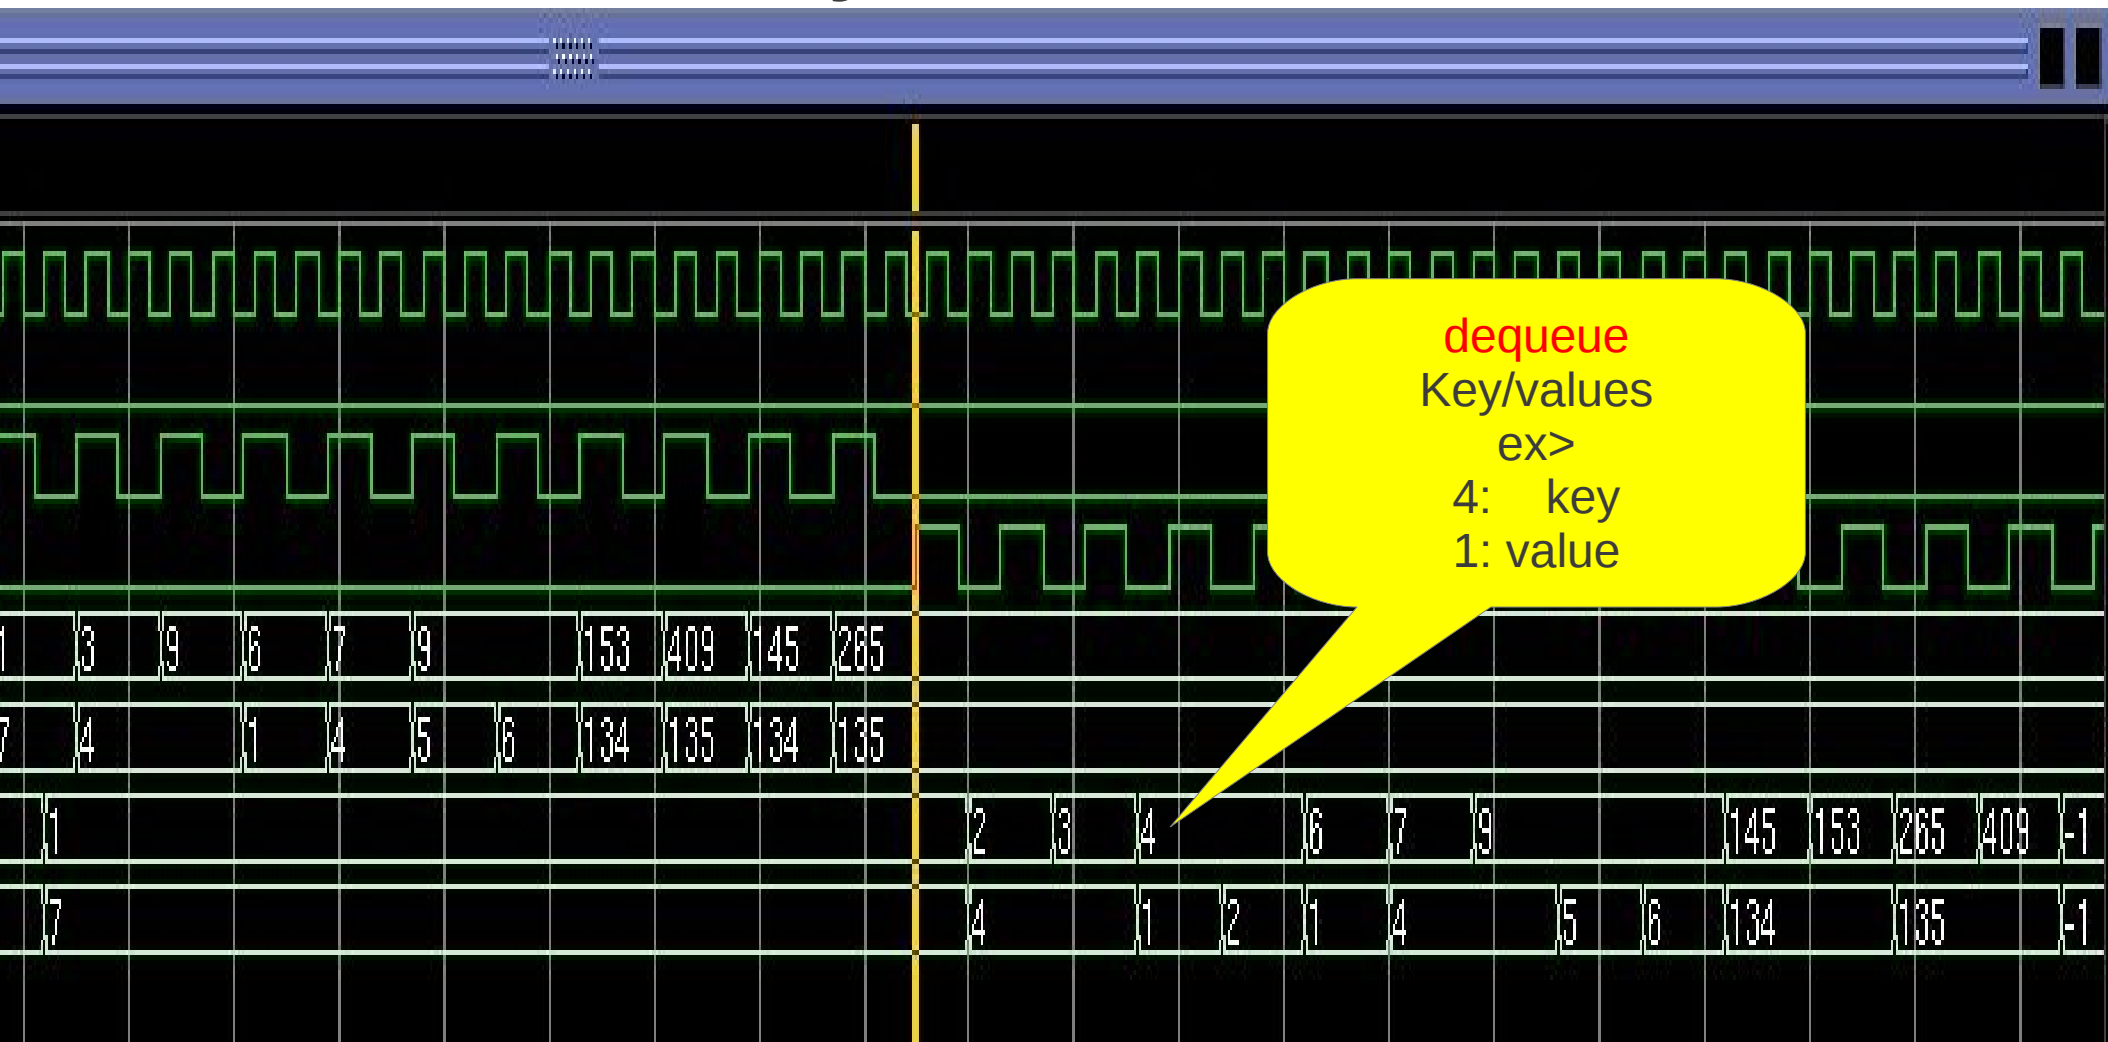


H/W Priority Queue Simulation





H/W Priority Queue Simulation



H/W Priority Queue Simulation



Scheduler change, `sched_fair.c`

```
/*
 * modified __enqueue_entity() to use Tagged Sorter
 * instead of rb-tree
 */
static void __enqueue_entity(struct cfs_rq *cfs_rq, struct sched_entity *se)
{
    struct rb_node **link = &cfs_rq->tasks_timeline.rb_node;
    struct rb_node *parent = NULL;
    struct sched_entity *entry;

    s64 key = entity_key(cfs_rq, se);
    __enqueue_tagged_sorter(key, se); // enqueue to Tagged Sorter
}
```

Scheduler change, `sched_fair.c`

```
/*  
 * modified __enqueue_entity() to use Tagged Sorter  
 * instead of rb-tree  
 */  
static void __enqueue_entity(struct cfs_rq *cfs_rq, struct sched_entity *se)  
{  
    struct rb_node **link = &cfs_rq->tasks_timeline.rb_node;  
    struct rb_node *parent = NULL;  
    struct sched_entity *entry;  
  
    s64 key = entity_key(cfs_rq, se);  
    __enqueue_tagged_sorter(key, se); // enqueue to Tagged Sorter  
}
```


Calling `h/w`
priority queue

H/W interface

```
/*  
 * enqueue entity to Tagged Sorter (HyOS)  
 */  
static void __enqueue_tagged_sorter(unsigned int key, unsigned int data)  
{  
    if (init == 0) {  
        hyos = (char*)ioremap(XPAR_HYOS_PLB_0_BASEADDR, 0x10000);  
        init = 1;  
    }  
  
    if (hyos != NULL) {  
        iowrite32(key,hyos+4*4); // key  
        iowrite32(data,hyos+5*4); // val  
        iowrite32(0xffffffff,hyos); // enqueue trigger  
        iowrite32(0x0,hyos);  
    }  
}
```

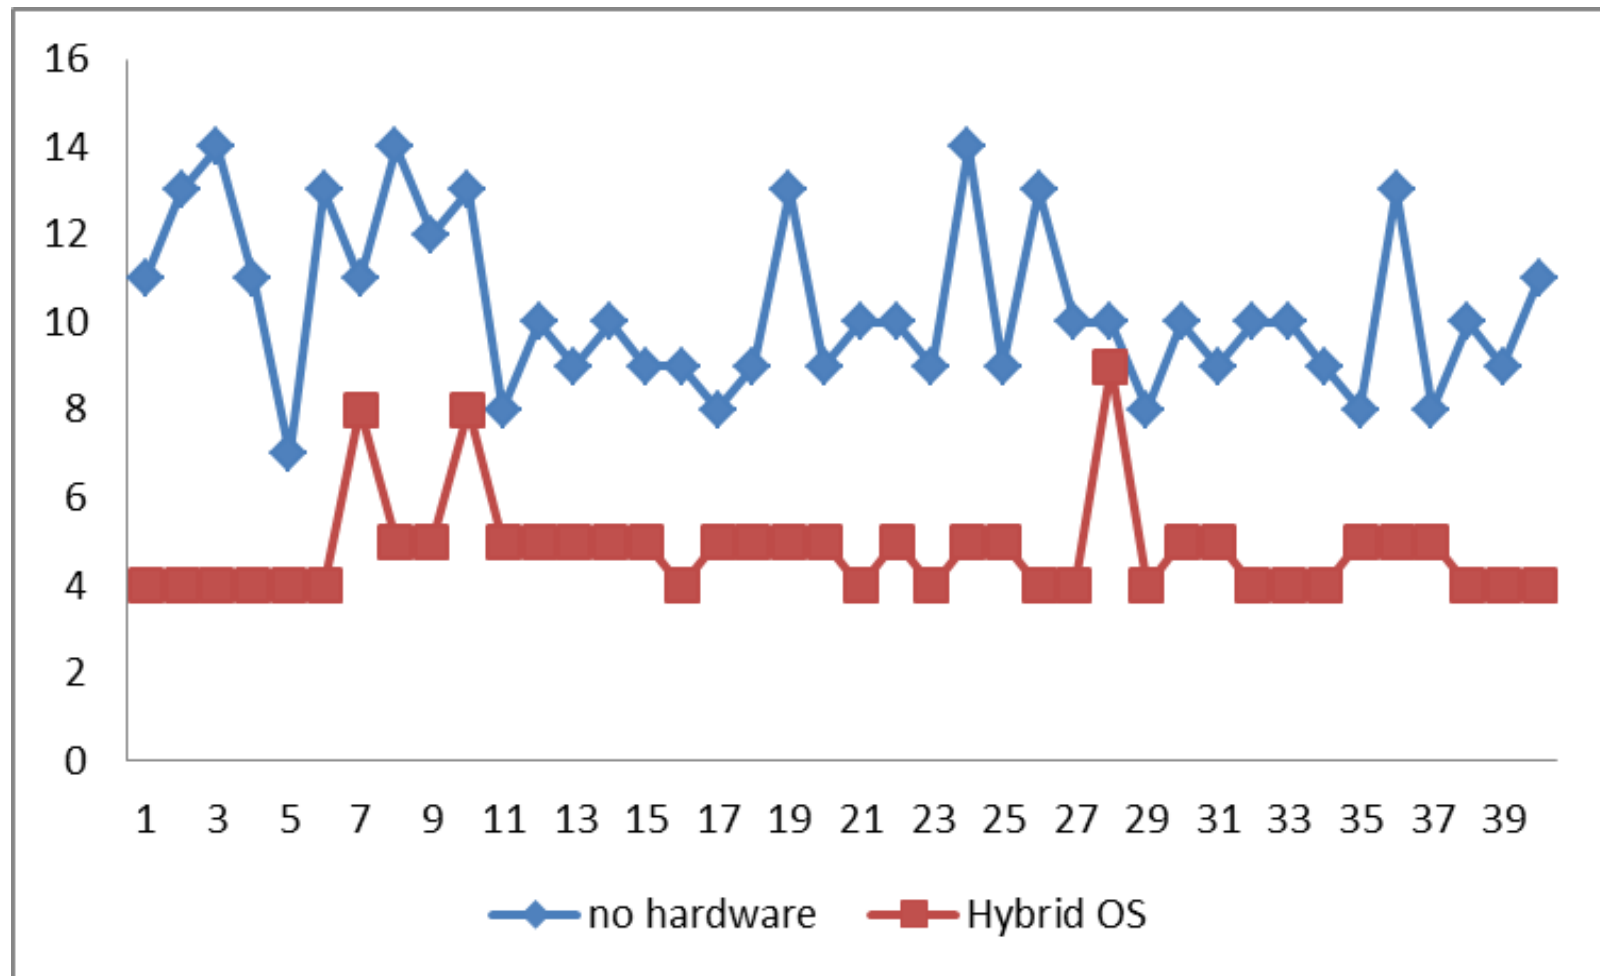
Experimental results 1

- Enqueue time for no load




Experimental results 2

- Enqueue time for 4 yes




Experimental results 2

- Enqueue time for 4 yes



Experimental results 2

- Enqueue time for 4 yes



S/W queue
increased &
more jitters

H/W queue
Stays the same

Conclusion

- H/W OS and H/W accelerator can
 - Improve performance
 - Reduce overhead
- Hybrid approach
 - Instead of implement full scheduler
 - Off-load a commonly used s/w component to h/w
- Could look for more s/w components could be off-loaded

Thank you!

Questions?